


Straightforward residual diagnostics for glm/glmer using the package DHARMa

Lisa Hülsmann

Standard residuals of generalized models are hard to interpret


```
plot(yourGLM)
```

- distributions are asymmetric
- residuals are never visually homogenous – even if the model is perfect
- more difficult for hierarchical models (random effects)
- but diagnostics are equally important
 - overdispersion (no variance parameter)
 - zero-inflation



Solution: Use the DHARMA package that is based on simulated residuals


- package developed by Florian Hartig
- user-friendly residual diagnostics for a large variety of models
- `simulateResiduals()`
 - generates residuals via simulation from the model
- `plotResiduals()`
 - plot residuals against predictors
- several test functions e.g.
 - `testDispersion()`
 - `testZeroInflation()`
 - ...



How to interpret the resulting plots


```
simulationOutput <- simulateResiduals(fittedModel = fittedModel)  
plot(simulationOutput)
```

DHARMA scaled residual plots






A misspecified model: here overdispersion

DHARMA scaled residual plots




Identify sources of mismatch: plot residuals against predictors

```
plotResiduals(YOURPREDICTOR, simulationOutput$scaledResiduals)
```



Several more advanced features

- calculating residuals per group
- temporal or spatial autocorrelation
- checking different levels of stochasticity in hierarchical models
- DHARMA can also be used for `lm/lmer`, `glmmTMB`, `smaMM` or posterior predictive simulations from Bayesian software



Check the vignette of DHARMA for more details.

You can also find nice examples for typical misspecifications there.

<https://cran.r-project.org/web/packages/DHARMA/vignettes/DHARMA.html>

Exercise

Data: abundance of myFavouriteSpecies
in response to temperature and precipitation,
with several observations made on a number of plots

count	temperature	precipitation	plot
2	-0.13159197	-0.031600868	1
1	0.15144222	-0.958592258	1
1	-0.10013614	-0.401583887	1
0	-0.97284735	-0.977506335	1
0	0.63306416	-0.647717534	1

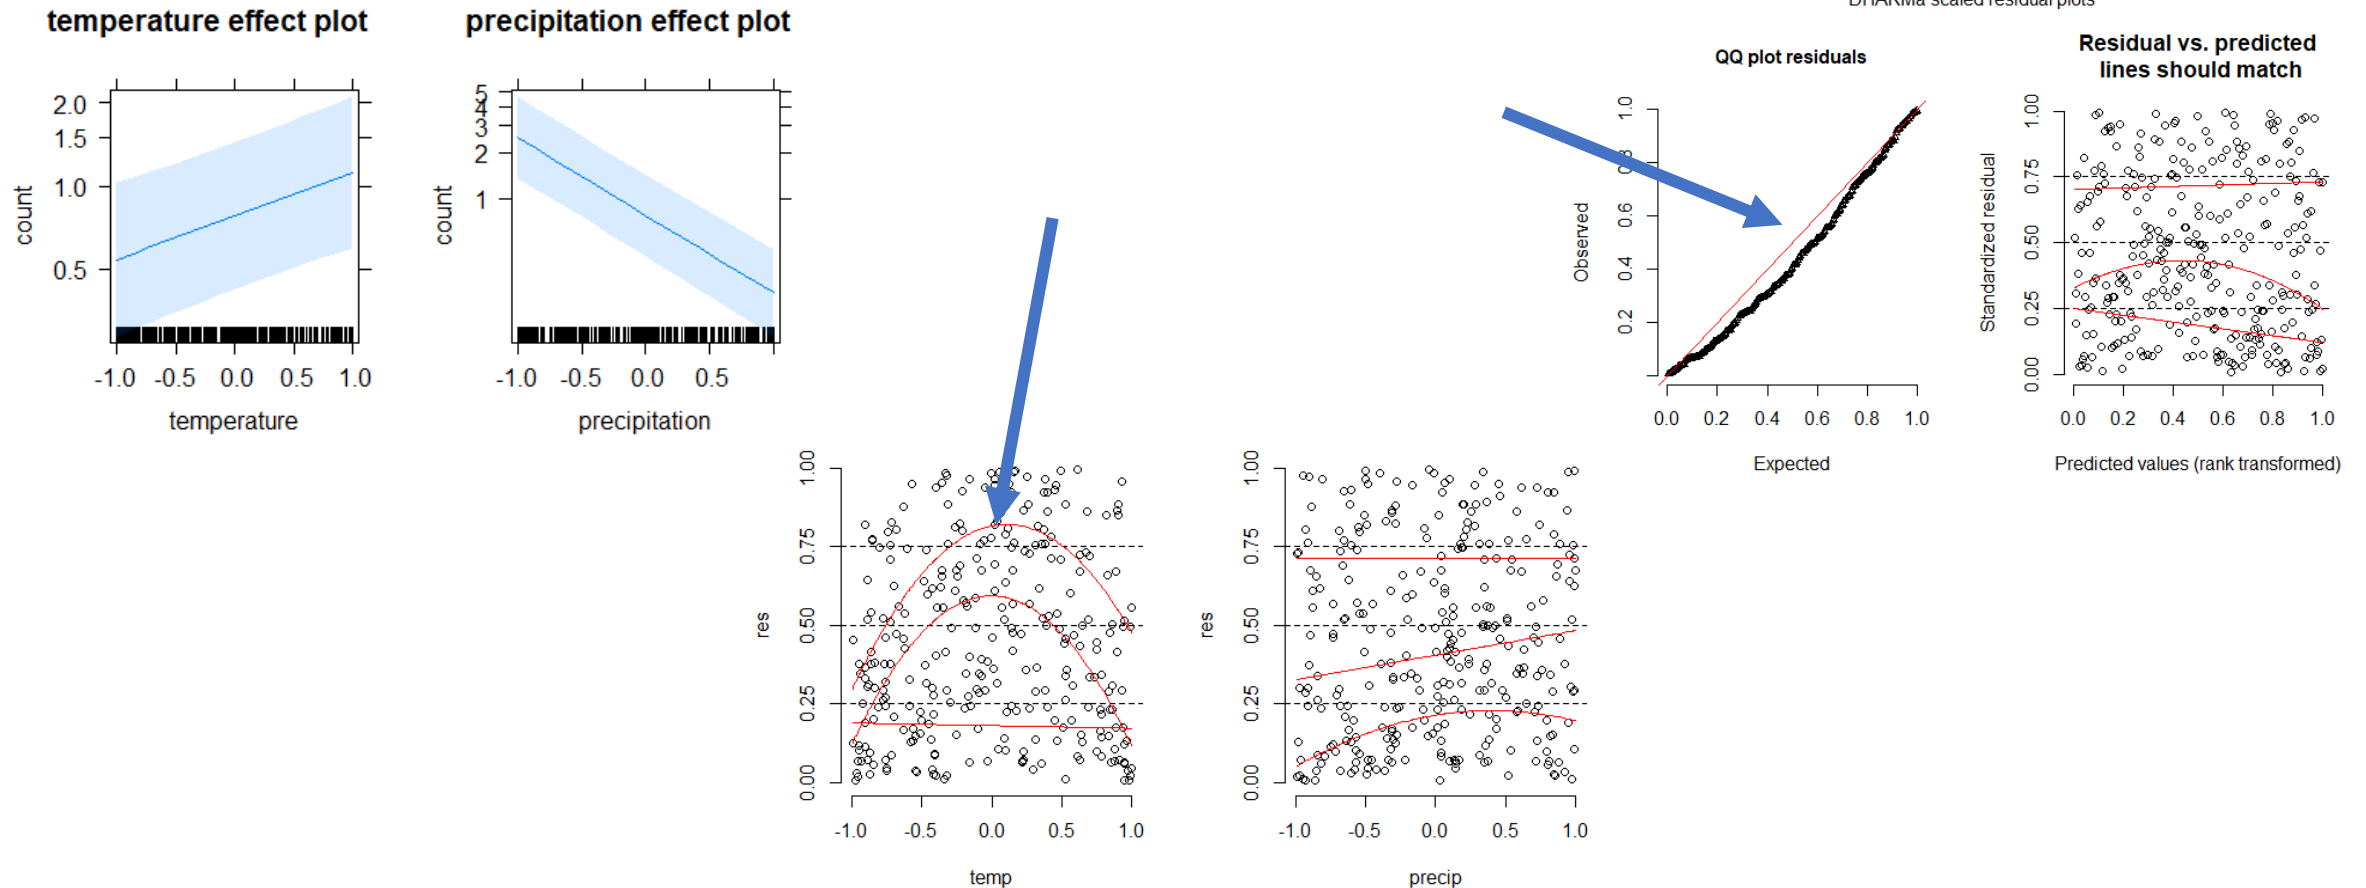
Task:

- fit a glmm with both predictors and use plot as a random effect (family = poisson)
- use the following workflow to check the model

```
simulationOutput <- simulateResiduals(YOURMODEL)
plot(simulationOutput)
plotResiduals(YOURPREDICTOR, simulationOutput$scaledResiduals)
```
- also try out the functions `testDispersion()` and `testZeroInflation()`
- what is wrong in the model? correct it and run the tests again
- How are species abundances influenced by temperature and precipitation

First model attempt


```
fittedModel <- glmer(count ~ temperature + precipitation + (1|plot),  
  family = "poisson", data = myFavouriteSpecies)
```






Clearly improved model

```
fittedModel <- glmer(count ~ temperature + I(temperature^2) + precipitation + (1|plot),  
                      family = "poisson", data = myFavouriteSpecies)
```


temperature effect plot



precipitation effect plot



QQ plot residuals



Residual vs. predicted lines should match

